

COMP 175

**System Administration
and Security**

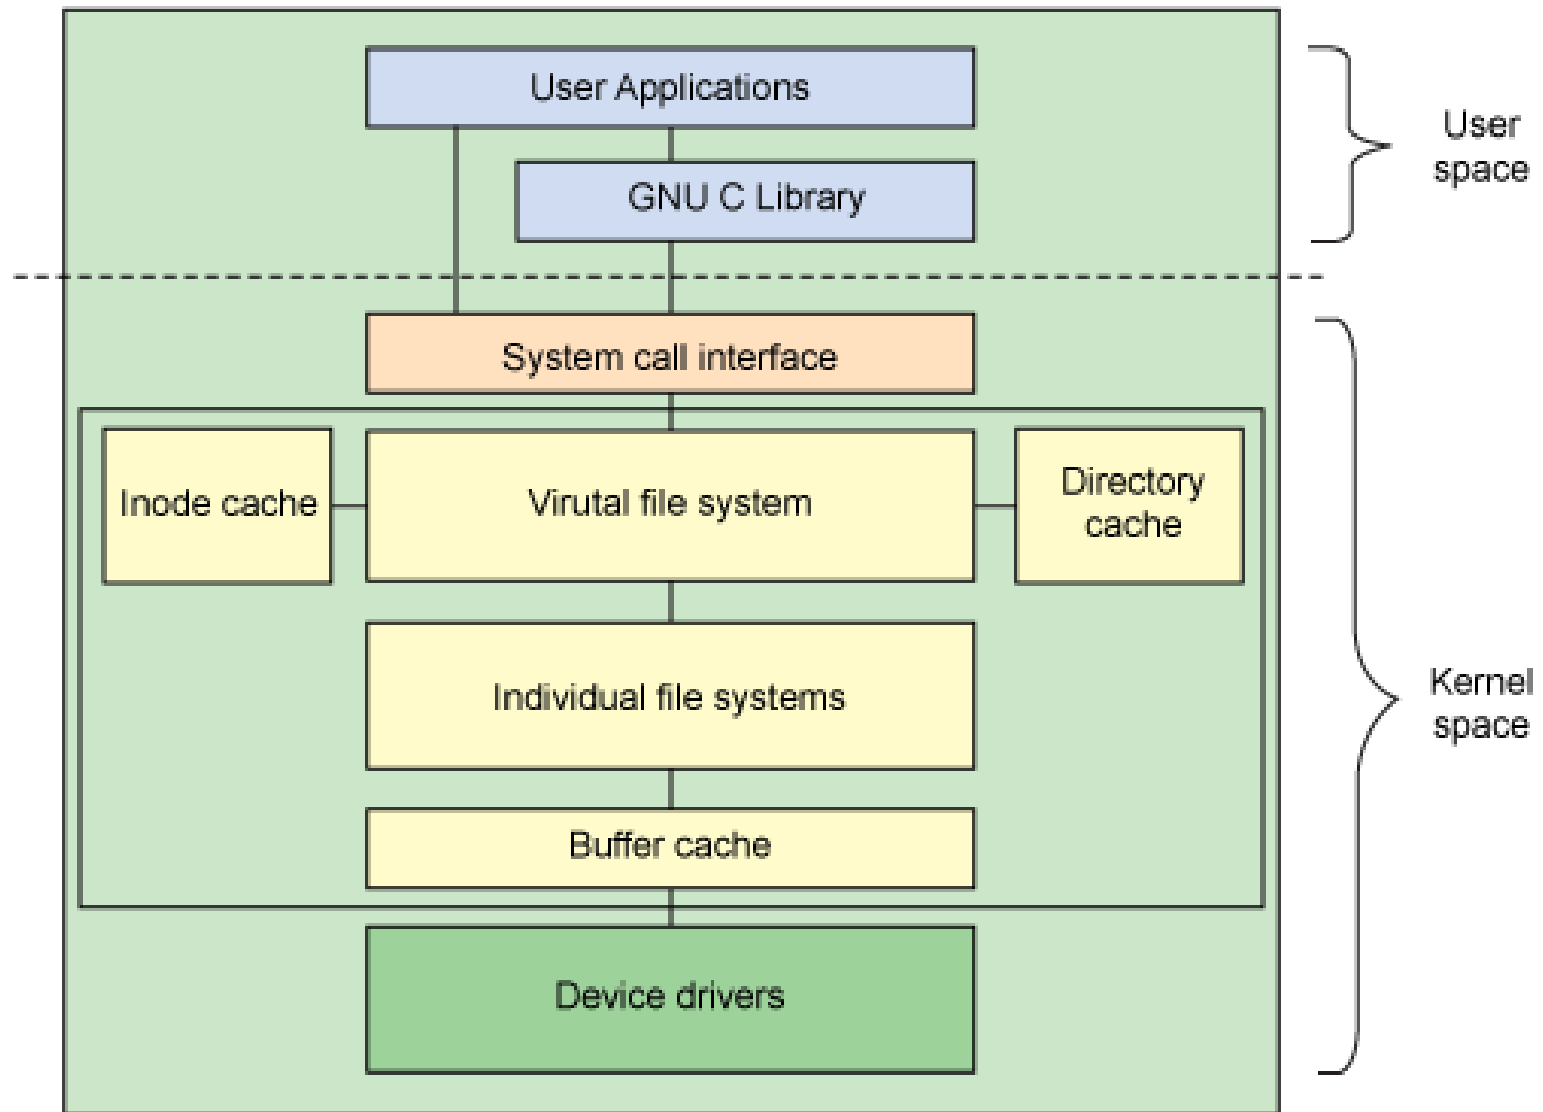


The File System

001000001 001000001 01010010 01010010 01010010



File System Architecture



Structural view of the file system

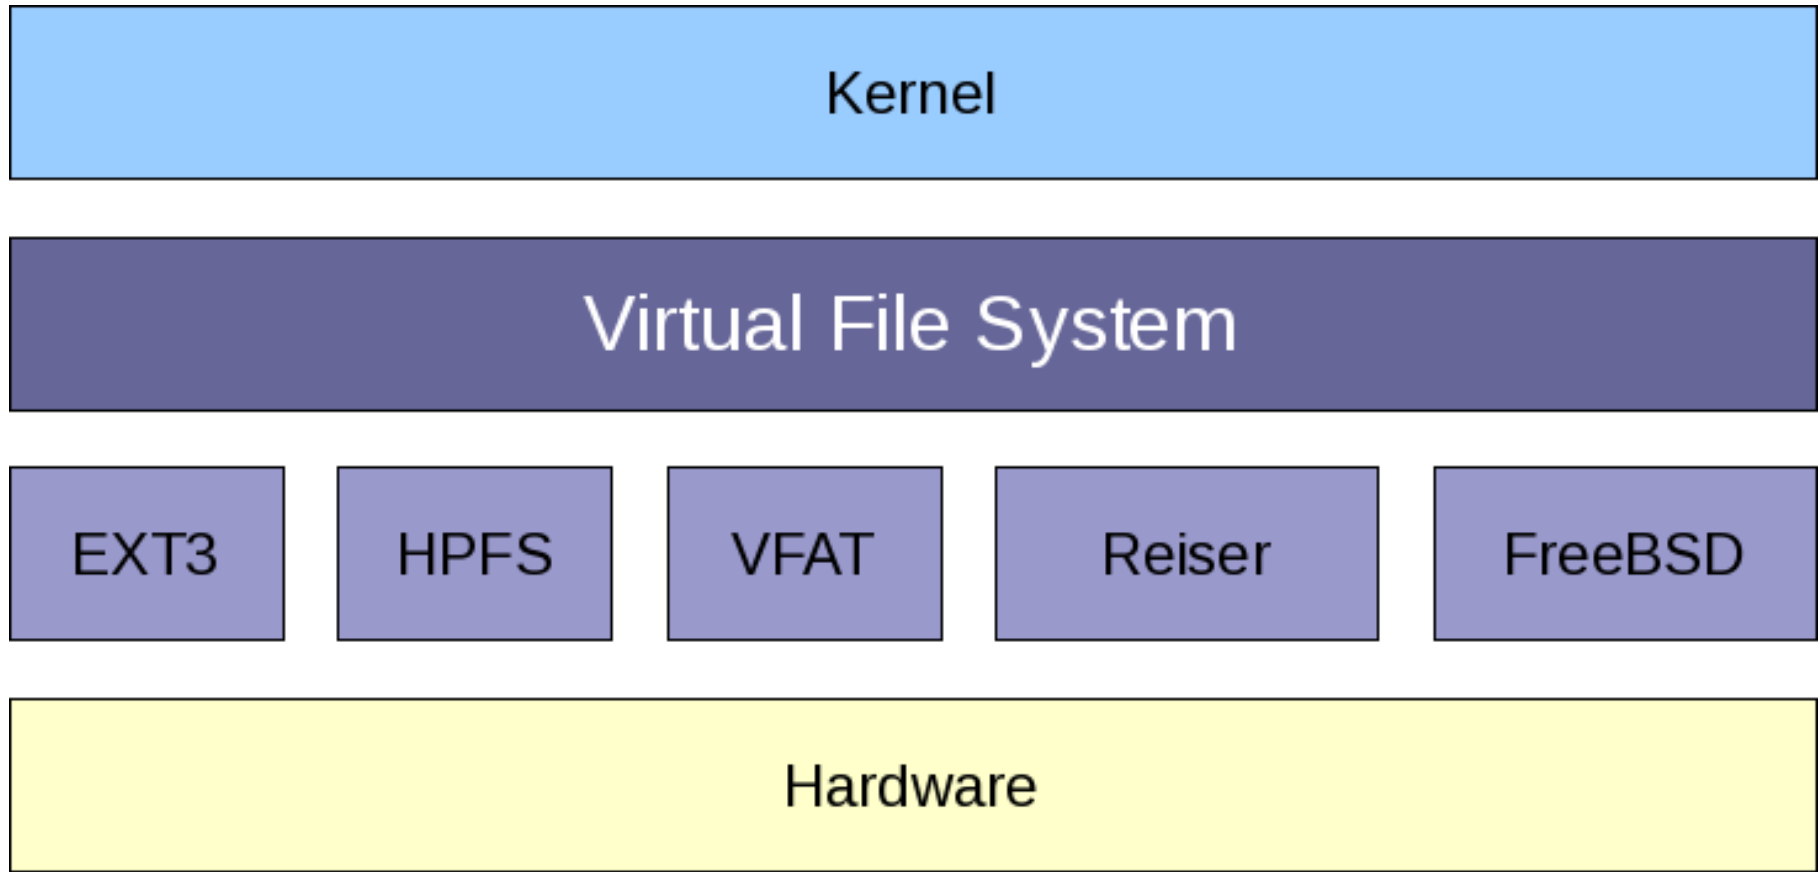


File System Architecture

- User space contains the applications (the user of the file system) and the GNU C Library (glibc), which provides the user interface for the file system calls (open, read, write, close). The system call interface acts as a switch, funneling system calls from user space to the appropriate endpoints in kernel space.
- The VFS is the primary interface to the underlying file systems. This component exports a set of interfaces and then abstracts them to the individual file systems, which may behave very differently from one another.
- Two caches exist for file system objects (**inodes** and **dentries** – directory entries). Each provides a pool of recently-used file system objects.



File System Architecture



Simplified further it looks like the above

With ~100 file systems supported – many legacy systems
one way to support legacy file systems



Virtual File System Layer

- VFS acts as the root level of the file system interface. The VFS keeps track of the currently-supported file systems, as well as those file systems that are currently mounted
- File systems can be dynamically added/removed
- Kernel keeps list of currently-supported file systems, which can be viewed from user space through the `/proc` file system.
- `/proc` virtual file also shows the devices currently associated with the file systems

We will get back to `/proc`



Buffer Cache

Buffer cache keeps track of r/w requests from the:

- Individual file system implementations
- Physical devices (through device drivers)
 - ◆ Linux maintains a cache of the requests
 - ◆ Avoids going back to physical device for all requests - efficient
 - ◆ The most-recently used buffers (pages) are cached thus can be quickly provided back to the individual file systems
 - Not found in cache? – Page fault



Page Faults

- The thread experiencing the page fault is put into a Wait state while the operating system finds the specific page on disk and restores it to physical memory
 - ◆ Takes time
 - Sends trap to the OS
 - Save user registers and process state
 - Determine location of the page on the disk
- Large numbers of page faults are an indication of insufficient RAM
- Also cause page reads (see disk counters)



Page Faults

- `top` add fault stats via `f nMaj` (not as flags)

```
top - 12:35:56 up 268 days, 1:17, 1 user, load average: 0.05, 0.04, 0.05
Tasks: 111 total, 1 running, 110 sleeping, 0 stopped, 0 zombie
Cpu(s):  0.0%us,  0.3%sy,  0.0%ni, 99.7%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Mem:    987208k total,  825080k used,  162128k free,  313392k buffers
Swap:   1694852k total,    0k used,  1694852k free,  358740k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	nFLT	COMMAND
2012	root	20	0	15660	3280	2496	S	0.0	0.3	0:00.14	14	smbd
1735	root	20	0	66508	36m	2680	S	0.0	3.8	52:31.06	12	named
1837	haldaemon	20	0	15576	4992	4232	S	0.0	0.5	4:04.56	10	hald
1765	root	20	0	25560	2568	2032	S	0.0	0.3	0:00.03	6	console-kit-dae
1	root	20	0	824	276	236	S	0.0	0.0	4:06.07	5	init
1836	root	20	0	21524	2648	2224	S	0.0	0.3	0:00.05	4	polkitd
2014	root	20	0	8620	2024	1396	S	0.0	0.2	23:53.24	3	nmdbd
1443	nobody	20	0	52672	1140	912	S	0.0	0.1	0:00.11	1	in.identd
1951	root	20	0	3772	1028	884	S	0.0	0.1	15:40.27	1	hald-addon-stor
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	0	kthreadd



Block & Character Devices

Block devices move data to/from that occur in blocks (such as disk sectors)

Supports buffering and random access behavior (is not required to read blocks sequentially, but can access any block at any time). Block devices include hard drives, CD-ROMs, RAM disks.

Character devices differ in that they do not have a physically-addressable media. Character devices include serial ports and tape devices, in which data is streamed character by character.




Journaling

- Journaling keeps track of the major steps taken during last file sessions
- If system crashes, it can boot and back up to the last known good configuration and recover to the point of the crash
- Hard drives may have their own write caches, journaling thus forces the device to flush its cache at certain points in the journal (called barriers in ext3 and ext4)
- ACID ([atomicity](#), [consistency](#), [isolation](#), [durability](#))
a set of properties that guarantee database transactions are processed reliably



Vocabulary

- Atomicity (database systems): a property of database transactions which are guaranteed to either completely occur, or have no effects.
- Atomicity states that database modifications must follow an “all or nothing” rule. Each transaction is said to be “atomic.” If one part of the transaction fails, the entire transaction fails. It is critical that the database management system maintain the atomic nature of transactions in spite of any DBMS, operating system or hardware failure





Linux File Systems

Linux has several standard file systems

- ext2 - legacy, general purpose, based on UFS
- ext3 – journaling, based on UFS
 - ◆ avoids long file system checks after system crash
- ext4 – journaling successor to ext2
- ReiserFS – journaling file system (infamous)
- XFS – IRIX's file system (good for streaming media)
- ZFS – Sun's combined file system & volume mgr
- procfs – interface to internal kernel structures
- swap – used to support virtual memory



Additional File Systems

Several foreign file systems are supported

- Easier to exchange files with another OS
- Work just like native ones, except:
 - ◆ May lack some usual UNIX feature
 - Long File Name support
 - UNIX permissions
 - ◆ Have curious limitations/oddities



CDROM file systems supported

- isofs – iso9660 CDROM file system
- Joliet - Microsoft CDROM filesystem extensions
 - ◆ Why is it called Joliet, you ask?



Joliet ISO 9660 Extension

- At the time of Windows NT 4.0 and Windows 95/98, the file system of choice to record files with names of 128 characters was called Romeo.
- Romeo didn't contain the ISO 9660 file system and broke backward compatibility with DOS
- Joliet combined Romeo and ISO 9660
- Joliet.doc metadata file in MS Win95 Driver Development Kit comment: "Joliet is a small town just outside of Chicago, where a man named Jake did some time in The Blues Brothers."



Elwood J. Blues, "Joliet" Jake B. Blues, Bluesmobile



Additional File Systems


- NFS - Network file system allows multiple users or hosts to share the same files using a client/server methodology
- NTFS – preferred Microsoft file system since NT provides ACL's and journaling
- FAT, FAT32 – potential 8.3 vs. Long Filenames (LFN), permissions complications
 - ◆ Used to mount usb flash drives
 - ◆ vfat – LFN compatible





Partitioning


- Partitioning is dividing a single hard drive into several logical drives.
- A partition is a contiguous set of blocks on a drive, treated as an independent disk.
- A partition table is an index that relates sections of the hard drive to partitions.





Multiple Partitions


- Multiple partitions reduce risk of system failure should a partition becomes full
- Segregating OS and user data space protects operating system if allocated disk space is exhausted
- Partitions can contain different **operating systems**
- Partitions can contain different **file systems**
- Called a *slice* in BSD, Solaris, and GNU Hurd





Physical Disk Structure

- 1 or more circular platters
- Platter has upper & lower oxide-coated surface
- Heads - min 1 per surface
 - ◆ Mounted on arms
- Heads float very close to platter surfaces
 - ◆ never touching them
 - ◆ disk crash





Partition Fields

- Device: the partition's device name
 - Start: drive sector where partition begins
 - End: drive sector where partition ends
 - Size: partition's size (in MB)
 - Type: partition type (e.g. ext2, ext3, or vfat)
 - Mount Point: where partition will be mounted within directory hierarchy (e.g. /, /var, /usr)
-
- Tools - fdisk, cfdisk
 - Always document settings – hard copy




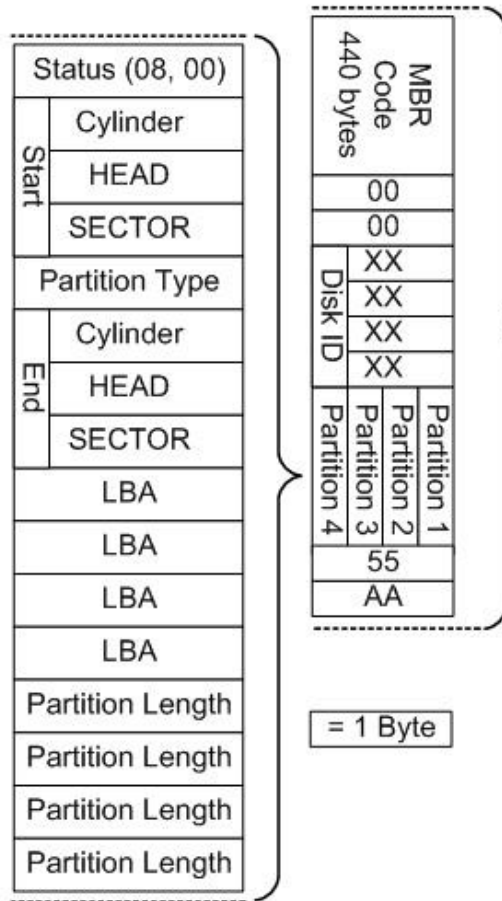
Master Boot Record

- MBR – 512 byte sequence at first sector of drive
- MBR used for one or more of:
 - ◆ Holds a partition (thus called a partition sector)
 - ◆ Bootstrapping an operating system.
 - PC BIOS loads the MBR from disk and passes execution to machine code instructions at the beginning of the MBR
 - ◆ Identify disk with a 32-bit disk signature
- MSDOS `fdisk /mbr` rewrites MBR – undocumented
- GRUB and LILO can write to the MBR



Master Boot Record

MBR in more detail



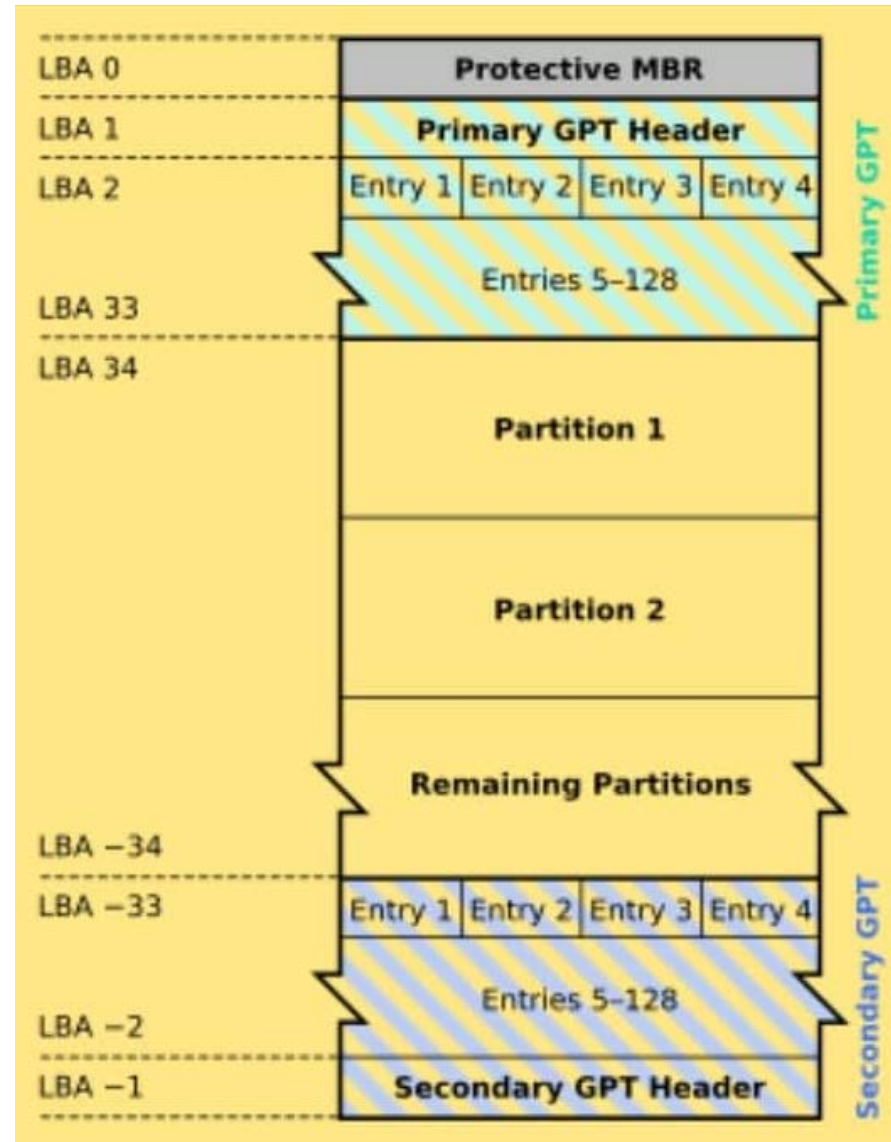
The partition definition section of the MBR could put them anywhere, they do not need to be outer to innermost





Master Boot Record

- The organization of the partition table in the MBR limits the maximum addressable storage space of a partitioned disk to 2 TiB ($2^{32} \times 512$ bytes).
- MBR-based partitioning scheme is in the process of being superseded by the GUID (Globally Unique Identifiers) Partition Table (GPT). Can co-exist with an MBR.
- GPT is part of UEFI standard





cfdisk

cfdisk 2.12

Disk Drive: /dev/hda

Size: 40020664320 bytes, 40.0 GB

Heads: 255 Sectors per Track: 63 Cylinders: 4865

Name	Flags	Part Type	FS Type	[Label]	Size (MB)
hda1	Boot	Primary	W95 FAT32		5239.51
hda2		Primary	Linux swap		403.04
hda3		Primary	Linux ReiserFS		10001.95
hda5		Logical	Linux ReiserFS		5996.23
hda6		Logical	Linux ReiserFS		5502.72
hda7		Logical	Linux ReiserFS		5502.72
hda8		Logical	Linux		5502.72
hda9		Logical	Linux		1867.14

[Bootable] [Delete] [Help] [Maximize] [Print]
[Quit] [Type] [Units] [Write]

Toggle bootable flag of the current partition



Disk Partition

IDE Disk Partition Example

Note: Physical drives can contain max. 4 primary partitions

- /dev/hda (Primary Master Disk)
 - ◆ /dev/hda1 (First Primary Partition)
 - ◆ /dev/hda2 (Second Primary Partition)
- /dev/hdb (Primary Slave Partition)
 - ◆ /dev/hdb1
- /dev/hdc (Secondary Master/Slave Partition)
 - ◆ /dev/hdc1



Disk Partition

SATA and SCSI Disk Partition Example

Note: Physical drives can contain max. 4 primary partitions

- /dev/sda (Primary Master Disk)
 - ◆ /dev/sda1 (First Primary Partition)
 - ◆ /dev/sda2 (Second Primary Partition)
- /dev/sdb (Primary Slave Partition)
 - ◆ /dev/sdb1
- /dev/sdc (Secondary Master/Slave Partition)
 - ◆ /dev/sdc1



Warning

- MultiBoot with Windows & Linux can be tricky
- Partitioning can be complicated
- Be careful altering partitions
 - ◆ Assume it will go bad
- Backup data prior to making any changes
- Read the documentation
- Research what others have done
- Check with other students
- Have a backout plan





Linux Disk Utilities

- `fdisk /dev/hda` Linux/DOS drive partitioning tool
- `cfdisk /dev/hda` Easier to use
- `sfdisk -l` Lists the partition tables
- `parted /dev/hda` Partition manipulation tool
- `fsck -t ext2 /dev/hda2` Check & repair file system

- `fsck` runs automatically at boot if OS detects files system wasn't properly shut down. Run when file system is unmounted or mounted read-only.
- Similar to MS scandisk or chkdsk



sfdisk

```
root@tea: ~# sfdisk -l
```

```
Disk /dev/hda: 1027 cylinders, 255 heads, 63 sectors/track
```

```
Units = cylinders of 8225280 bytes, blocks of 1024 bytes, counting from 0
```

Device	Boot	Start	End	#cyls	#blocks	Id	System
/dev/hda1	*	0+	995	996-	8000338+	83	Linux
/dev/hda2		996	1026	31	249007+	82	Linux swap
/dev/hda3		0	-	0	0	0	Empty
/dev/hda4		0	-	0	0	0	Empty



sfdisk

- `sudo sfdisk -l` (edited) Ubuntu
 - ◆ gpt
 - ◆ UEFI

```
Disk /dev/sda: 698.64 GiB, 750156374016 bytes, 1465149168 sectors
Disk model: ST9750423AS
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 4096 bytes
I/O size (minimum/optimal): 4096 bytes / 4096 bytes
Disklabel type: gpt
Disk identifier: 959BD32A-80AA-4293-900E-A1C42542517F
```

Device	Start	End	Sectors	Size	Type
/dev/sda1	2048	4095	2048	1M	BIOS boot
/dev/sda2	4096	1054719	1050624	513M	EFI System
/dev/sda3	1054720	1465147391	1464092672	698.1G	Linux filesystem



Inconsistent State

- A non-graceful **shutdown** (crash, power loss) can leave a file system in an Inconsistent State. Prior to journaling file systems, it was common for an improperly shut-down Unix system's file system to develop a corrupted superblock.
- Running fsck to fix this could take minutes to hours (volume size and disk I/O throughput)
- The consequences of fsck not being able to fix the error are not good
- Hence: "**fsck**" and "fscked"



fsck on OS-X

Note: Spelling auto-correct doesn't play well with UNIX terms



```
Run /sbin/fsck -y first and then /sbin/mount -uw
sh-2.05a# fsck -f
** /dev/rdisk0s9
** Root file system
** Checking HFS Plus volume.
** Checking Extents Overflow file.
** Checking Catalog file.
Overlapped extent allocation (file 9447270d)
** Checking multi-linked files.
Orphaned indirect node temp9117989
Orphaned indirect node temp9291654
Orphaned indirect node temp9447245
Orphaned indirect node temp9447270
Orphaned indirect node temp9447294
Orphaned indirect node temp9447309
** Checking Catalog hierarchy.
** Checking volume bitmap.
Volume Bit Map needs minor repair
** Checking volume information.
Invalid volume free block count
(It should be 5173737 instead of 5179225)
** Repairing volume.

***** FILE SYSTEM WAS MODIFIED *****
sh-2.05a#
```




Remember

- MBR: Master Boot Record
- Superblock is a boot record
- Big-endian machine stores most significant byte first
- Little-endian machine stores least significant byte first

Ponder this

- Journaling file systems
- Raid
- What if you don't have ECC?
- Error correction code memory (ECC memory) uses an error correction code (ECC) to detect and correct n-bit data corruption which occurs in memory. ECC memory is used in most computers where data corruption cannot be tolerated, like industrial control applications, critical databases, and infrastructural memory caches.